

THE ANTHROPIC BRIEFING

Eleven days that took Anthropic from a new flagship model to the doorstep of the public markets

WEEKLY EDITION · VOL. 1 NO. 1

COVERAGE 25 MAY – 4 JUN 2026

COMPILED FROM 7 DAILY BRIEFINGS

This issue covers a uniquely dense stretch. In a single afternoon — **28 May** — Anthropic shipped **Claude Opus 4.8** and disclosed a **\$65B Series H** at a **\$965B post-money valuation**. Four days later it **confidentially filed a draft S-1 with the SEC**, turning a year of IPO rumors into a documented option. Underneath the headlines, Claude Code shipped **ten releases**, the agent SDKs realigned, Project Glasswing tripled in size, and the reliability log filled with model incidents — capped by a major Claude Code outage on 3 June. Here is the whole arc, distilled for people who build on it.

Opus 4.8 NEW FLAGSHIP, SHIPPED 28 MAY — SAME \$5/\$25 PRICE AS 4.7	\$965B SERIES H POST-MONEY VALUATION · \$65B RAISED	S-1 DRAFT FILED CONFIDENTIALLY WITH THE SEC, 1 JUN
10× CLAUDE CODE RELEASES, V2.1.152 → V2.1.162	~150 PROJECT GLASSWING PARTNERS, UP FROM ~50, IN 15+ COUNTRIES	33 → 56% CYBER ACTORS RATED MEDIUM-RISK- OR-HIGHER, H1 → H2

IN THIS ISSUE

- | | |
|--|--|
| 01 Lead — A model, a mega-round, an IPO path | 06 Research & engineering |
| 02 Models & the API | 07 Reliability watch |
| 03 Claude Code — ten releases in eleven days | 08 Rumors, leaks & money |
| 04 The SDK lockstep | 09 From the community |
| 05 Security & cyber — Glasswing, Mythos, ATT&CK | 10 The practitioner's checklist |

01 A model, a mega-round, and a path to the public markets

LEAD STORY · THE THROUGHLINE OF THE WINDOW

If you read only one thing about this period, read it as three moves in one direction. On **28 May**, Anthropic made Claude Opus 4.8 generally available *and* announced a **\$65 billion Series H** that put post-money valuation at **\$965 billion**, with self-reported run-rate revenue that "crossed \$47B earlier this month." The round was led by Altimeter, Dragoneer, Greenoaks, and Sequoia, co-led by Capital Group, Coatue, D1, GIC, ICONIQ, and XN, and folded in \$15B of previously-committed hyperscaler money

(including \$5B from Amazon) plus memory/chip partners Micron, Samsung, and SK hynix. Several outlets framed it as Anthropic overtaking OpenAI as the most valuable AI company.

Then, on **1 June**, the corporate-structure shoe dropped: Anthropic **confidentially submitted a draft Form S-1 to the SEC**. It is a Rule 135 announcement — no share count, no price — so it legally separates "we *can* IPO" from "we *will*." But it is the move that typically precedes a 6–12-month IPO window, and it converts the long-running funding rumors (a "\$30B at \$900B" leak we tracked on 27 May; the WSJ/TechCrunch "first profitable quarter" narrative) into something concrete you can cite in a vendor-risk conversation.

The funding rumors are now first-party numbers. \$47B run-rate and a \$965B valuation are figures you can put in a budget memo — not press speculation.

SERIES H ANNOUNCEMENT · 28 MAY 2026

The skeptics arrived on schedule. Michael Burry ("The Big Short") publicly argued neither SpaceX nor Anthropic is worth \$1T (134 pts on Hacker News), and *The Economist's* "Can the stockmarket swallow Anthropic, SpaceX and OpenAI?" was the single most-discussed thread of the window at **712 points / 1,257 comments**. None of that changes the operational picture — but if you are quoting the Series H valuation to justify spend, expect the counter-narrative in the same meeting.

WHY IT MATTERS

The through-line for builders: the company you depend on is simultaneously shipping its most capable model at flat pricing, raising at a near-trillion-dollar valuation, and laying IPO groundwork. Read that as *stability of roadmap* — and as growing pressure to convert subsidized usage into per-token revenue, the pricing shift Simon Willison flagged this week (see \$9).

02 Models & the API

WHAT CHANGED FOR PEOPLE CALLING THE MESSAGES API

Claude Opus 4.8 — a same-price capability bump

Model string `claude-opus-4-8`, GA 28 May, available everywhere (API, Claude Code, claude.ai, Cowork) on day one. The platform facts that matter for building:

- **Pricing unchanged from 4.7:** \$5 / Mtok input, \$25 / Mtok output; fast mode \$10 / \$50. Fast mode on 4.8 is roughly **3× cheaper** than on prior models (2× the standard rate for 2.5× the speed).
- **1M-token context by default; 128K max output.** Effort defaults to `high` on every surface, including the raw API, with *adaptive thinking* (reasoning only fires when a turn needs it). `/effort xhigh` exists for the hardest tasks.
- **Prompt-cache minimum lowered to 1,024 tokens** — cheaper to cache short system prefixes than on 4.7.
- **Reliability claim:** ~4× **less likely** than Opus 4.7 to let flaws in its own code pass unremarked; the Alignment team reports misaligned-behavior rates "similar to" Mythos Preview. Launch-tester

benchmarks: 84% on Online-Mind2Web (computer use), first model past 10% overall on a Legal Agent Benchmark, Genie multimodal at 61% lower token cost than 4.7.

Two Opus-4.8-only Messages API features — no beta header

- **Mid-conversation system messages.** You can now place `role:"system"` entries at non-first positions in the `messages` array to update instructions mid-task *while preserving prompt-cache hits* — instead of routing changes through a user turn. Directly useful for long-running agents that need to revise permissions, token budgets, or environment context mid-run.
- **Refusal categories in `stop_details`.** When Opus 4.8 declines, the API returns a refusal category so you can route different refusal classes to the right next step instead of string-matching error text. By 3 June (Claude Code v2.1.162 / agent-SDK v0.3.162) the same `stop_reason:"refusal"` and `stop_details` also land on local-session and SDK transcripts.

MULTI-CLOUD NOTE

As of Claude Code v2.1.158 (30 May), **Auto mode is available on Bedrock, Vertex, and Foundry** for Opus 4.7 and 4.8 — opt-in via `CLAUDE_CODE_ENABLE_AUTO_MODE=1`. Previously the auto-approval classifier was first-party only. If you route through a hyperscaler for compliance or billing, you no longer give up the layer that (per Anthropic) catches ~83% of overeager behaviors.

No new model since. Opus 4.8 remained the newest first-party model through 4 June; Mythos stays gated to Project Glasswing (see §5). The base `anthropic-sdk-*` packages did not move in the final stretch — the Opus-4.8 feature adds landed in `anthropic-sdk-python` v0.105.0 and `anthropic-sdk-typescript` sdk-v0.100.0 on launch day.

03 Claude Code — ten releases in eleven days

THE RELEASE TRAIN, DECODED

The window opened on v2.1.150 and closed on v2.1.162. Below is the practitioner-relevant signal, release by release (v2.1.151/.155/.159 were internal-only and are omitted from detail).

VERSION	DATE	WHY YOU CARE
2.1.152	27 May	<code>/code-review --fix</code> writes findings back to your tree; <code>/simplify</code> now calls it. <code>disallowed-tools</code> skill/command frontmatter; <code>/reload-skills</code> + <code>SessionStart</code> → <code>reloadSkills</code> ; new <code>MessageDisplay</code> hook. Auto mode no longer needs opt-in consent. Fixes under-reported <code>cache_creation_input_tokens</code> (expect a metrics discontinuity).
2.1.153	28 May	Security: fixed a gateway credential leak (custom gateway could receive the user's own OAuth token) and subagent-frontmatter MCP servers bypassing <code>--strict-mcp-config</code> / managed allow-deny policy. <code>/model</code> now saves as the default; the picker keybinding <code>d</code> → <code>s</code> .
2.1.154	28 May	The Opus 4.8 vehicle. Ships dynamic workflows (research preview — plan + tens-to-hundreds of background subagents + self-verification; Enterprise/Team/Max). Lean system prompt becomes default; <code>/simplify</code> is cleanup-only. Improved auto-mode detection of bulk-exfil.
2.1.156	29 May	Hotfix: thinking blocks were being modified on Opus 4.8, causing API errors. If you run 4.8 in Claude Code, this is the floor.

VERSION	DATE	WHY YOU CARE
2.1.157	29 May	Plugins in <code>.claude/skills</code> auto-load, no marketplace ; <code>claude plugin init</code> scaffolds one. Per-session <code>agent</code> override (<code>--agent</code>). OTel <code>tool_decision</code> gains <code>tool_parameters</code> . Worktree-hygiene fixes.
2.1.158	30 May	Auto mode on Bedrock / Vertex / Foundry for Opus 4.7 & 4.8 (<code>CLAUDE_CODE_ENABLE_AUTO_MODE=1</code>).
2.1.160	2 Jun	Dynamic-workflow trigger renamed <code>workflow</code> → <code>ultracode</code> (asking in your own words still works). New write-confirmation prompts before touching shell startup files and exec-granting build configs (<code>.npmrc</code> , <code>.bazelrc</code> , <code>.pre-commit-config.yaml</code> ...). <code>grep</code> now satisfies read-before-edit. <code>CLAUDE_CODE_OPUS_4_6_FAST_MODE_OVERRIDE</code> removed.
2.1.161	2 Jun	<code>OTEL_RESOURCE_ATTRIBUTES</code> become metric labels (slice usage by team/repo). Parallel tool calls are now independent — one failed Bash call no longer cancels the batch. Secret-redaction in <code>claude mcp</code> list/get/add .
2.1.162	3 Jun	<code>claude agents --json</code> adds <code>waitFor</code> . <code>--tools</code> accepts <code>Grep</code> / <code>Glob</code> on native builds again . Autocomplete click now <i>fills</i> a slash command (press Enter to run). WebFetch deny/allow rules now override the preapproved-host allowlist. MCP <code>timeout</code> < 1000 ms no longer aborts every call.

HEADS-UP — SECRETS & THE `WORKFLOW` RENAME

Two items that bite if you miss them: **(1)** the `workflow` trigger keyword is gone — update any docs, scripts, or hooks that key off the word "workflow" to `ultracode` . **(2)** Secret-redaction in `claude mcp` only arrived in v2.1.161; treat any earlier scraped `claude mcp` output or shared session screencap as having leaked credential headers and URL secrets.

Catch-up from the gap: v2.1.149 (22 May) was the most material missed release — `/usage` per-category breakdown (skills, subagents, plugins, MCP cost) plus a cluster of sandbox-boundary security fixes (PowerShell `cd` bypass, worktree write-allowlist over-scoping).

04 The SDK lockstep

AGENT SDKS REALIGN AFTER A 10-DAY PYTHON STALL

For most of the window the Python agent SDK lagged a full Claude Code cycle behind TypeScript. That gap is now closed.

- `claude-agent-sdk-typescript` moved in lockstep with Claude Code (v0.3.152 → v0.3.162). Notable: v0.3.160 fixed hook callbacks swallowing abort signals; v0.3.161 made `initialize` idempotent and live-applies `agent` changes; v0.3.162 changed the **default search behavior on native builds** — SDK sessions now default to embedded `find` / `grep` in Bash rather than registering the dedicated `Grep` / `Glob` tools.
- `claude-agent-sdk-python` resumed shipping after stalling at v0.2.87 since 23 May: **v0.2.88** (2 Jun) ported session stores from raw `asyncio` to `anyio` , fixing a Trio crash; **v0.2.89** (3 Jun) is a CLI-parity bump. TS and Python are back in the same cohort (v0.3.162 / v0.2.89 / CLI 2.1.162).

ACTION ITEM

If you run an SDK orchestrator with **hook-based search auditing** (`canUseTool` / interception on `Grep` / `Glob`), verify it still fires after upgrading — the dedicated tools are no longer registered by default on native builds. Name them explicitly in `tools` / `allowedTools` to keep the old behavior.

05 Security & cyber — Glasswing, Mythos, and a year of threat data

THE MOST CONSEQUENTIAL RESEARCH THREAD OF THE WINDOW

Project Glasswing tripled — and pivoted to patching

The 2 June "Expanding Project Glasswing" post (178 pts / 247 comments) is the substantive update. Partners went from ~50 to ~150 organizations across 15+ countries, now spanning power, water, healthcare, communications, and hardware — well beyond the original software cohort. Partners have surfaced **more than 10,000 high-or-critical-severity flaws**, and the program made its first explicit pivot from *finding* vulnerabilities to *disclosing, fixing, and deploying patches* — with Mythos Preview now writing patches and running pre-release checks. **Claude Security** is named as the productized form (Opus 4.8 scanning codebases and suggesting patches).

READ THE FINE PRINT ON "MYTHOS IN THE COMING WEEKS"

The Glasswing post explicitly conditions a Mythos-class general release on "highly robust safeguards... that we (and, to our knowledge, all other AI developers) have *yet to develop*." Treat the 28 May launch line about Mythos arriving "in the coming weeks" as **conditional, not committed**. A new "Cyber Verification Program" would grant Mythos-class access to additional orgs for specific cyberdefense tasks — no public details yet.

A year of AI-enabled cyber threats, mapped to MITRE ATT&CK

The 3 June Frontier Red Team post analyzed **832 accounts** banned for malicious cyber activity (Mar 2025 – Mar 2026; partial results also published in Verizon's 2026 DBIR). The findings reset some common heuristics:

- **Threat level is inflating.** The share of actors rated medium-risk-or-higher rose from **33% in the first half to 56% in the second** (~1.7×).
- **AI is moving deeper into the kill chain.** AI-assisted phishing (initial access) fell 8.6%; AI for post-compromise account discovery rose 8.9%. Malware-writing is still the most common use (67.3%); lateral movement is rare (6.5%) but a bigger risk signal.
- **Old signals broke.** Technique count and which surface was used (Claude Code vs API vs chat) *no longer* correlate with risk. The durable signal is **scaffolding architecture** — whether the actor chains attack stages with autonomous decision-making. Anthropic is in talks with MITRE about extending ATT&CK with agentic-orchestration techniques.

IF YOU BUILD ABUSE / DLP / AGENT-AUDIT TOOLING

The "number of techniques" heuristic is now publicly deprecated by the vendor. Feature-engineer toward **whether model-generated activity chains stages autonomously**, not how many distinct techniques it touches.

The Mythos evidence trail

Across the window the Mythos story moved from rumor to hard data point. A `claude-mythos-1-preview` model string and a "Claude Security" dashboard were reported (unconfirmed) inside Claude Code on 25 May. More concretely, **Apple credited "Calif.io in collaboration with Claude and Anthropic Research"** for kernel integer-overflow bug **CVE-2026-28952** (macOS Tahoe 26.5) — and a high-voted HN thread (145 pts) surfaced an Anthropic engineer's on-record line that the team "built a working exploit in five days" with Mythos Preview. It is the strongest external confirmation yet of the Glasswing/Mythos workflow producing a shipped, CVE'd OS-vendor vulnerability.

06 Research & engineering

TWO READS WORTH YOUR TIME

"How we contain Claude across products" — the highest-leverage single read

Published 25 May, it front-paged on HN (139 pts) on 4 June and is now the canonical citation for the v2.1.160–162 design choices. The numbers practitioners keep quoting:

- Users approve roughly **93%** of permission prompts. Claude Code auto mode cuts prompts by **84%** and catches **~83%** of overeager behaviors before they execute...
- ...but auto mode blocks only **~0.4%** of benign commands while **~17%** of overeager actions still get through — so it is *defense-in-depth*, not a substitute for a sandbox.
- Three containment tiers documented (gVisor for claude.ai; Seatbelt/bubblewrap HITL sandbox for Claude Code, with the runtime open-sourced; full local VM for Cowork), plus three on-record near-miss incidents and Gray Swan numbers (Opus 4.7 holds attack success to **~0.1%** single-shot, **~5–6%** after 100 adaptive attempts).

"Coding agents in the social sciences" — adoption, quantified

A 27 May Economic Research post surveyed **1,260 quantitative social scientists**. The headline gap: 81% have tried AI chatbots, but only **20% regularly use coding agents** — and among those, **Claude Code dominates at 86%** (Codex 31%). Adoption is steeply uneven (economists 39%, public health 6%; >2× the rate for typically-male vs -female names; top-25 universities ~40% higher). Usage is overwhelmingly code, not prose. Rare first-party, quantified evidence for an internal adoption/ROI case — and a preview of the disparity findings that will surface in governance conversations.

"I think Anthropic and OpenAI have found product-market fit."

SIMON WILLISON · 27 MAY · 892 PTS / 1,011 COMMENTS ON HN

07 Reliability watch

A NOISY ELEVEN DAYS ON STATUS.CLAUDE.COM

The incident log was the busiest in recent memory — an Opus 4.7 drumbeat early, the first *major* Opus 4.8 incident two days post-launch, and a multi-hour Claude Code outage to close.

25 May	MAJOR	Opus 4.7 elevated errors, ~4 hours (06:30–10:30 UTC) — second 4.7 incident in five days.
27 May	MAJOR	Two back-to-back Opus 4.7 incidents (~58 min + ~1h37m) across all four surfaces.
28 May	MAJOR	Opus 4.7 elevated errors, ~39 min — the fourth 4.7 major in roughly a week.
29 May	MAJOR	First major Opus 4.8 incident (18:35–19:12 UTC), ~2.5 days after launch.
2 Jun	MAJOR	Elevated errors on multiple models, ~5h45m.
3 Jun	MAJOR	"Issue affecting some Claude Code services," ~3h20m (04:17–07:36 UTC) — the biggest Claude-Code-specific incident of the window.

ACTION ITEM

Four major Opus 4.7 events in a week, plus a major 4.8 event and a multi-model outage, make **retry-with-fallback wiring non-optional** for production traffic in the EU/Asia-morning UTC window. Note the limit: Claude Code's "switch to `--fallback-model` when the primary is *not found*" does **not** cover 5xx / elevated-error events on a model that is found — you still need your own backoff + fallback path.

08 Rumors, leaks & money

WHAT RESOLVED, WHAT SOFTENED, WHAT'S STILL AIR

- **Funding — RESOLVED.** The "\$30B at \$900B" leak (27 May) and the WSJ/TechCrunch revenue rumors were superseded by the official **\$65B Series H at \$965B** (28 May). Cite the first-party numbers, not the leaks.
- **IPO — CONFIRMED as an option.** The draft S-1 (1 Jun) makes the IPO path first-party. Valuation price-discovery hasn't happened, so 11-figure valuations remain projections.
- **"Opus 4.8 today?" — RESOLVED true.** The single-X-post rumor from 28 May shipped that afternoon.
- **"Mythos in the coming weeks" — SOFTENED.** The 2 Jun Glasswing post conditions GA on safeguards no AI developer has yet built (see §5).
- **Counter-narrative.** Michael Burry's sub-\$1T skepticism is the highest-vote skeptic thread — useful framing to have ready, not an operational signal.

09 From the community

PRESS, TOOLING, AND THE ENTERPRISE-ADOPTION DEBATE

The PMF essay. Simon Willison's "product-market fit" piece (892 pts) argues enterprise coding agents are the labs' PMF moment, now that both Anthropic and OpenAI have moved enterprise plans to API-rate pricing. His own `ccusage` run: ~\$1,200/mo of Claude Code tokens for \$200 in subscription — the "17× subsidy" framing that also ran as its own Show HN.

The skeptical counterweight. Fortune reported **Uber burned its entire 2026 AI coding-tools budget in four months**, with its COO unable to "draw a line" between Claude Code usage stats and shipped customer value (CEO: ~10% of committed code is agent-built). The same cycle carried The Verge's report that **Microsoft began canceling most of its direct Claude Code licenses**, moving engineers to GitHub Copilot CLI — part cost, part dogfooding. The retention question to be ready for: tool-use stats vs. shipped value.

Interpretability, in public. Co-founder Chris Olah's remarks at the Vatican (on Pope Leo XIV's AI encyclical) included an on-record attestation that the interp team has found "internal states that functionally mirror joy, satisfaction, fear, grief, and unease." No paper yet — a forward signal more than a citable result.

Tooling & field notes. A "Claude Code as a Daily Driver" guide (CLAUDE.md, skills, subagents, plugins, MCPs) drew an active HN thread; "Everything You Can Configure That the Docs Don't Tell You" (121 pts) is a power-user reference. *Launch HN: Hyper (YC P26)* signals a wave of "company memory layer" startups targeting the Claude Code / Cowork hook surface. On the org side, Anthropic appointed KiYoung Choi as Korea Representative Director ahead of a Seoul office.

10 The practitioner's checklist

TEN DAYS OF CHANGES, DISTILLED TO WHAT YOU SHOULD ACTUALLY DO

1. **Rename your triggers.** Update docs, scripts, and hooks that key off the word `workflow` → `ultracode` (v2.1.160).
2. **Floor your Claude Code version at 2.1.156+** if you run Opus 4.8 (thinking-block API-error fix); 2.1.162 is current.
3. **Re-test hook-based search auditing.** Native SDK builds now default to embedded `find` / `grep`; name `Grep` / `Glob` in `tools` / `allowedTools` if you intercept them.
4. **Add retry-with-fallback** for model incidents — `--fallback-model` won't save you from 5xx on a model that exists.
5. **Rotate anything exposed.** Treat pre-v2.1.161 `claude mcp` output / screenshots as leaked secrets.
6. **Enable Auto mode on Bedrock/Vertex/Foundry** (`CLAUDE_CODE_ENABLE_AUTO_MODE=1`) if you route through a hyperscaler.
7. **Check stray MCP `timeout` values.** A `timeout:500` was silently aborting every tool call before v2.1.162; sub-1000 ms now falls back to the default.

- 8. Reach for mid-conversation system messages for long-running agents that need to revise instructions without busting the prompt cache (Opus 4.8 only).
- 9. Expect a cache-metrics discontinuity at v2.1.152 — `cache_creation_input_tokens` was previously under-reported in some configs.
- 10. Re-tool your abuse/agent-audit features toward autonomous stage-chaining, not technique counts (3 Jun MITRE post).

. . .

§ Sources & method

IN-WINDOW, FIRST-PARTY-VERIFIED WHERE POSSIBLE

First-party — Introducing Claude Opus 4.8 + System Card · Series H announcement · Confidential draft S-1 (Rule 135) · Expanding Project Glasswing · "What we learned mapping a year of AI-enabled cyber threats" (MITRE ATT&CK) · "How we contain Claude across products" (eng) · "Coding agents in the social sciences" (research) · Chris Olah / Pope Leo XIV encyclical remarks · KiYoung Choi appointment · platform & support release notes.

Releases — Claude Code v2.1.149 → v2.1.162 (GitHub releases + CHANGELOG) · `claude-agent-sdk-typescript` v0.3.152 → v0.3.162 · `claude-agent-sdk-python` v0.2.88 / v0.2.89 · `anthropic-sdk-python` v0.105.x · `anthropic-sdk-typescript` `sdk-v0.100.x`.

Status & community — `status.claude.com` incidents API (3 major + multiple minor in window) · Hacker News via Algolia `search_by_date` · The Economist, NYT, Fortune, The Verge, Nature, Apple security advisory (CVE-2026-28952), Simon Willison.

Opus 4.8 · 28 May

Series H · \$965B

Draft S-1 · 1 Jun

Glasswing ×3

MITRE ATT&CK

v2.1.152→162

workflow→ultracode

CVE-2026-28952

6 major incidents

Compiled from seven daily Anthropic briefings spanning 2026-05-25 (morning & afternoon editions), 05-26, 05-27, 05-28, 05-29, and the 06-04 catch-up briefing (which covered the 29 May – 4 Jun gap). Figures and quotations are reproduced as stated in those briefings; where a source flagged a number as unverified or a page date as unconfirmed, that caveat carries forward. Dates are UTC. Generated 2026-06-06 for an applied-AI-engineering readership.